

A Framework for Preprocessor-Aware C Source Code Analyses

Greg J. Badros and David Notkin

Technical Report UW-CSE-98-08-04

Dept. Computer Science & Engineering

University of Washington

Box 352350, Seattle WA 98195-2350 USA

+1-206-543-1695

{gjb,notkin}@cs.washington.edu

ABSTRACT

Analyses of C source code usually ignore the C preprocessor because of its complexity. Instead, these analyses either define their own approximate parser (or scanner) or else they require that their input already be preprocessed. Neither approach is entirely satisfactory: the first gives up accuracy (or incurs large implementation costs), while the second loses the preprocessor-based abstractions. We describe a framework that permits analyses to be expressed in terms of both preprocessing and parsing actions, allowing the implementer to focus on the analysis. We discuss an implementation of such a framework that embeds a C preprocessor, a parser, and a Perl interpreter for the action “hooks.” Many common software engineering analyses can be written surprisingly easily using our implementation, replacing numerous ad-hoc tools. The framework’s integration of the preprocessor and the parser further enables some analyses that otherwise would be especially difficult.

Keywords

Parsing, preprocessor, `cpp`, lexical analysis, syntactic analysis, Perl.

1 INTRODUCTION

More than twenty years ago, Dennis Ritchie designed the C language [11] to include a textual macro preprocessor called `cpp` [10, Ch. 3]. Given the simplicity of the language and the state of the art in compiler technology in the mid-1970s, his decision to provide some language features in this extra-linguistic tool was justified. For the last two decades, C programs have exploited `cpp`’s capabilities for everything from manifest constants and type-less pseudo-inline functions to modularization and symbol generation. Bjarne Stroustrup, the designer and original implementor of C++, notes that “without the C preprocessor, C itself . . . would have been stillborn” [21, p. 119]. Certainly `cpp` contributes to C’s expressiveness and portability, but perhaps at too large a cost. Stroustrup recognizes this tradeoff:

Occasionally, even the most extreme uses of `cpp` are useful, but its facilities are so unstructured and intrusive that they are a constant problem to programmers, maintainers, people porting code, and tool builders [21, p. 424].

Why `cpp` is good . . . *and* bad

The intrinsic problem with `cpp` is also its fundamental strength: it is a distinct first pass of textual (non-syntactic) processing over the source code. This introduces significant differences between the code that the programmer sees and what the compiler proper (i.e., the C compiler distinct from the preprocessor) ultimately compiles.¹

Experienced and novice C programmers alike are frustrated by misunderstandings of source code due to the arbitrary transformations `cpp` performs. Such confusions are easily reduced, if not eliminated by allowing the software engineer to see the code exactly as the compiler does. Unfortunately, that view of the program is at a level of abstraction lower than the unprocessed source. Well-known identifiers such as `stderr` may appear as the far less readable `(_IO_FILE*)(&_IO_stderr_)`, and useful encapsulations such as `assert` degenerate into sequences of symbols that are less meaningful to a human programmer. Every non-trivial C program uses the preprocessor, and an empirical study of numerous packages written in C documents the extensive use of `cpp` constructs [2].

`cpp` and Software Tools

Because of the preprocessor’s textual foundations, unprocessed C source code cannot be parsed directly. For example, identifiers which are macro-expanded may hide braces, parentheses, or other important tokens from the parser. Only after preprocessing is a C program in a grammatically usable form. Since parsing pre-

¹To avoid ambiguity, we will use *unprocessed* to refer to the original source code, and *preprocessed* to refer to the view of the code after running `cpp` on it.

processed code is relatively easy and well-understood, most software engineering tools operate on exactly that view of the source, losing abstractions that are expressed in `cpp`.

Various tools including source-level debuggers and call graph extractors either run `cpp` as the first stage in their analysis, or use representations derived from a compiler operating on the preprocessed code. Siff and Reps implement a function-generalization transformation tool that operates on preprocessed code, but note that future versions of their tool must drop that requirement so that macro abstractions are not lost [19]. Using preprocessed code for program understanding or transformation tasks is fraught with difficulties due to changes in the source artifact from preprocessing.²

Another disadvantage of preprocessing is that it eliminates conditional compilation directives that are essential to the portability and versatility of the source code [12]. Preprocessing forces tools to limit their analysis to a single configuration of the source code, instead of permitting global reasoning about the entire artifact.

Some tools instead operate on the unprocessed source code exactly as the programmer sees it. This technique improves robustness in handling syntax errors and language variants. Because the input is unprocessed, the extracted information is presented to the human software engineer at the same level of abstraction as the source. Additionally, the unprocessed source still contains the preprocessor directives which are essential to the portability and flexibility of many C programs.

However, these tools cannot use a straightforward parser or reliably construct an abstract syntax tree because the source is not grammatically-correct C code. Lexical tools such as `etags`, `LSME` [15], and `font-lock-mode` for Emacs and approximate parsers such as `Field` [18] and `LCLint` [3] use this approach. Unfortunately, disregarding or only partially honoring the `cpp` directives leads to an extracted model of the source code that is only an approximation of the program’s appearance to a compiler. Certain uses of the preprocessor can cause substantial problems: the syntax of declarations or scoping constructs can be customized and arbitrary code can be hidden in macro expansions. Macro expansion was a major cause of both false positives and false negatives in call graph extraction [9]. Such approximate tools are inappropriate for software engineering analyses that require exact or conservative information.

²In contrast, this approach is exactly right for the compiler, which has no need to preserve high level abstractions while generating object code.

We introduce a new approach that integrates the C preprocessor and a parser in a flexible framework for statically analyzing C source code. The framework makes it easy for tool builders to produce analysis tools for C source code that are capable of reasoning about the preprocessor. Section 2 describes that framework and illustrates a sample analysis (a call-graph extraction) using the framework. Section 3 details a pair of macro-expansion understanding analyses, and section 4 describes our prototype tool, `PCp`³. Finally, section 5 discusses related work, and sections 6 and 7 review the limitations of our approach and conclude.

2 THE FRAMEWORK

The idea behind our framework is simple: provide an integrated parser and preprocessor which controls the scanning of the source code while executing user-defined callbacks when specified “interesting” actions occur. Though similar to the way the `yacc` parser [13] associates actions with parse rule invocation, our framework provides callbacks on both parser and preprocessor actions. Additionally, instead of C as the language for the actions, we use the Perl scripting language for writing the hook subroutines. Griswold, Atkinson and McCurdy note that various software tools benefited from using a special-purpose action language [9], and interpreted languages can often significantly speed development time [17].

Callbacks can be installed on actions such as the scanning of preprocessor directives, the creation of a macro definition, the expansion (i.e., use) of a macro name, and the parsing of a variable declaration. Each action callback is passed arguments relevant to the event for which it was invoked. For example, the `TYPDEF` action receives the name of the declared type as its first argument. See Appendix A for a more complete list of the hooks in our current implementation of the framework.

To supplement the directly-passed arguments, action subroutines may also invoke “backcalls” to access internal parser and preprocessor data structures. Example backcalls include getting the name of the currently-processed file, looking up a symbol in the parser’s symbol table, inserting arbitrary code for the preprocessor to process, and instructing the parser to enter a new scope. See Appendix B for a more complete list of the backcalls in our current implementation of the framework.

The callback and backcall interfaces combined with a general purpose scripting language provide a concise and flexible means for easily writing static analyses of C source code.

```

use pcp3;
my %func_calls = (); # the dictionary of callees for the current function
                      # maps called function name to number of times it appears
                      # in the definition of the calling function

AddHook("FUNC_SPEC", sub { %func_calls = (); } );
AddHook("FUNC_CALL", sub { # $_[0] argument is the name of function invoked
                          $func_calls{$_[0]}++; } );
AddHook("FUNCTION",
  sub { # $_[0] argument is the name of function just defined
    if (scalar (keys %func_calls) > 0) {
      print "$_[0] calls ", join(", ", sort keys %func_calls), "\n";
    }
  } );

```

Figure 1: A complete static analysis to extract a call graph.

Example: Call-Graph Extraction

One common software engineering analysis is to extract a call graph from a body of C source code. Our framework permits this analysis to be written in about ten lines of code, as shown in Figure 1.

To perform the analysis, we attach subroutines to three parsing callbacks. When a function specifier (the signature specification of a function definition) is parsed, the `FUNC_SPEC` hook is activated. That callback resets the set of functions invoked by the current function. For each function call parsed, the `FUNC_CALL` hook is activated, and we add the name (passed to the subroutine as the zeroth argument) of the called function to the set. Finally, when the entire function definition has been parsed, the `FUNCTION` reduction occurs and we output the set of called functions we accumulated while scanning the body of the function definition.

The script does not implement the main control flow of the analysis. Also, it need not handle any preprocessor directives. Since this simple analysis deals only with parser actions, the view the analysis sees is exactly the same as if we had analyzed the preprocessed code.

Revising the Extractor

The simple call-graph extractor used as an example in the preceding section suffers from many of the same shortcomings as tools which operate on preprocessed code: it reasons about only a single configuration of the source and eliminates all macro abstractions from the extracted model. As Murphy et al. discuss, there are many degrees of freedom for call-graph extractors [16]. A strength of our framework is that it lets a tool-builder easily fine-tune the extraction to derive the desired view.

For example, often macros are used to express inline functions. If we choose, we can have macro expansions included in the extracted call-graph:

```

# $_[2] is the macro name being expanded
AddHook("EXPAND_MACRO",
  sub { $func_calls{$_[2]}++; }

```

Given this revised extractor, a macro that expands into code that includes function calls will expose those function calls as callees from the function definition in which the macro was expanded. For some tasks, this may be exactly what we want. If instead we prefer to hide those nested calls, our framework supports that behaviour as well. We simply ignore function calls that we parse while expanding macros (see Figure 2).

Another possible extension of the analysis is to query the symbol table for the type of the identifier being called. If the variable is a pointer to a function, the enclosing function could conservatively be marked as “possibly calling all functions” (i.e., \perp).³

Because the activations of the preprocessor and parser hooks are intermingled, the analysis can reason about the preprocessed version of the source within the context of how it was preprocessed.

Preprocessor and Parser Interactions

Given a straightforward preprocessor and parser front end, some analyses are impossible. For example, since the preprocessor will skip code as instructed by con-

³A more aggressive analysis could maintain data on which functions are ever assigned to the variable throughout the program.

```

use pcp3;
my %func_calls = (); # the set of callees for the current function
my $expanding = 0;   # are we expanding a macro?

AddHook("FUNC_SPEC", sub { %func_calls = (); });
AddHook("FUNC_CALL", sub { # $_[0] argument is the name of function invoked
                           $func_calls{$_[0]}++ if !$expanding; } );
AddHook("EXPAND_MACRO", sub { # $_[2] argument is macro name being expanded
                              $func_calls{$_[2]}++ if !$expanding;
                              ++$expanding } );
AddHook("MACRO_CLEANUP", sub { --$expanding; } );

AddHook("FUNCTION",
  sub { # $_[0] argument is the name of function just defined
    if (scalar (keys %func_calls) > 0) {
      print "$_[0] calls ", join(", ", sort keys %func_calls), "\n";
    }
  } );

```

Figure 2: A revised call graph extractor that treats macros as functions.

ditional compilation directives, that source will be left unprocessed.

To permit handling the otherwise-skipped code hidden by conditional compilation directions such as `#ifdef` and `#if`, our framework provides a general mechanism for inserting arbitrary source text to be processed. Action subroutines can pass a string to the `PushBuffer` backcall. A hook, `DO_XIFDEF`, is called for all conditional compilation directives and provides an argument that is the text that would *not* be included. The action routine for that hook may then use `PushBuffer` to ask the preprocessor and parser to use that string as input and process it normally.

However, since such program text can include arbitrary preprocessor directives and C code, we may want to ensure that the state of neither the preprocessor nor the parser is permanently changed after the “for analysis only” parsing of the skipped code. Three data structures must be preserved to avoid side-effects to the main processing (i.e., the version selected by the `#definedness` of macro names): the preprocessor’s hash table of macros, the C symbol table, and the parser’s current stack of states.

The backcalls `PushHashTab` and `PopHashTab` save and restore the preprocessor’s table of definitions so that preprocessor directives in skipped code will not affect `cpp` when normal processing resumes. For the symbol table, the backcalls `EnterScope` and `ExitScope` provide similar functionality. Finally, `YYPushStackState` and

`YYPopStackState` save and restore the stack of parser states. In under fifty lines of boiler-plate code, an analysis tool can manage these complexities and expose the full source to the analysis.

Other related backcalls provide additional support for querying and interacting with the parser. `ParseStateStack` returns the list of states in the parser’s stack; this exposes information about what constructs might be legal at the current location (for example, determining whether declarations may be permitted). `SetStateStack` permits explicitly changing the parser’s stack of states, perhaps to prepare for handling a top level construct or to re-parse text using a different configuration of the parser. Using `YYCompareTopStackState`, an action hook can efficiently check whether a previously saved parser state matches the current configuration. For example, an analysis can easily provide a warning whenever the “if” block and the “else” block of an `#ifdef` directive leave the parser in different states.

3 FURTHER EXAMPLES

Preprocessor-specific analyses are generally especially difficult to write. The complexities and subtleties of `cpp` must be duplicated in each tool. Not surprisingly, our framework is ideal for analyzing preprocessor constructs.

Marking Macro Expansions

For an empirical study of C preprocessor use [2], we wanted to identify macro expansions in the unprocessed

source code. Recognizing identifiers that are macro-expanded is exceedingly useful as suggested by the nearly universal convention of using all-uppercase identifiers for macro names. Previous tools have difficulty even just counting macro expansions [9]. For the empirical study, our first approximation of identifying macro expansions did not use our framework. That analysis was overly conservative: we marked all occurrences of any identifier that was ever (in any source file of the program) `#defined` to be a macro as an expansion. To gain more accurate information, we used our framework for the subsequent version of the analysis.

Since the marking of expansions should cover the entire unprocessed source code, we use the boiler-plate code mentioned in the previous section to expose all the code to the tool. Because we wanted a conservative analysis (i.e., it should over-mark when imprecise), we treat an identifier as a macro at a program point if it has been `#defined` on any path and not yet definitely `#undefined`. In particular, the improved analysis properly limits the effects of macros that are defined for a segment of code and subsequently undefined.

Our improved analysis using the framework is written in under 90 lines of code, most of which is the boiler-plate for handling all conditionally-included code.

cpp Transformation Understanding

When debugging preprocessor difficulties, experienced C programmers have learned the hackish usefulness of directly viewing the output of `cpp`. A more precise character-by-character mapping between the unprocessed and preprocessed source proves even more beneficial to understanding the transformation performed by `cpp`. Our framework permits deriving such a fine-grained mapping with relative ease (about 300 lines of Perl callbacks and code).

Since the analysis is meant to be used as an interactive program-understanding tool, we needed a useful way to present the extensive information about the mapping. To achieve this, the tool outputs Emacs Lisp code [14] which marks up the unprocessed source using text properties. A custom Emacs mode then permits viewing of those annotations (see Figure 3).

Other Possibilities

Our framework has proven useful for several software engineering analyses. Other ad-hoc tools could benefit from our framework, as well. For example, Emacs’s `etags` [14] creates a database of function definitions in unprocessed source but can be easily confused by package-specific macros used in function definition

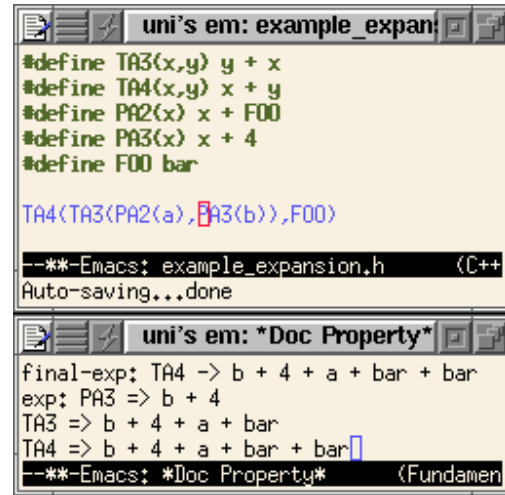


Figure 3: An interactive view of the mapping between unprocessed and preprocessed views of an example C program fragment. The outlined box in the top frame indicates the user’s cursor position and the lower frame describes the macro-expansion information at that character location. The bottom frame potentially changes after each cursor movement.

headers.⁴ Also, searching for uses of global variables could be enhanced using our framework since it will not omit references to globals hidden by macros.

Another possible use of the framework is in identifying “tricky” uses of the preprocessor as an aid to determining, for example, which `#defines` can be replaced with language-proper features such as constant variables or C++ inline functions.

4 THE PCp³ TOOL

Our framework combines a **P**arser, a **C** preprocessor, and a **P**erl action language. Thus, we have named the tool that implements the framework PCppP or PCp³.

Parser

Choosing a parser was difficult as there are many freely available parsers, often tightly coupled to a functional back-end, thus complicating reuse. We chose the parser from CTree [4], a freely available ANSI C front end, to embed in PCp³ largely because of its simple implementation and fully-scoped symbol table. Its lexical analyzer and parser are mechanically generated by `flex` and `bison` [5, 13] (freely available implementations of `lex` and `yacc`, respectively). As CTree parses, it builds a

⁴To circumvent this limitation, `etags` permits the programmer to manually specify an arbitrary regular expression which it uses to find definitions to mark.

complete abstract syntax tree of the preprocessed code.

The implementation of the **CTree** parser component of **PCp³** is about 5,000 lines of C code and **bison** and **flex** specifications. Most of the changes to the parser were to eliminate name conflicts and to call the Perl hooks as part of various reduction rules.

Preprocessor

Since **PCp³** must mimic **cpp** exactly, the C preprocessing library from the GNU C compiler's (**gcc**) well-tested (and slightly extended) **cpp** [6] is embedded in **PCp³**.

The implementation of **cpplib** grew from about 7,000 lines of code as distributed with **gcc** to almost 8,000 lines. Most of the changes involved modifying data structures and function calls to maintain the extra macro expansion information to support the character-by-character correspondence between the source and output.

Action Language

We chose Perl as the action language for **PCp³** because it interfaces easily with C and is used widely. Additionally, Perl's built-in data structures, closures, and higher-order functions make it especially well-suited for use with the framework.

The Perl subroutine hooks are written in a user-specifiable script file that registers a subroutine for each action it wants to process. That script is free to manage its local data structures arbitrarily, and may import reusable modules as an ordinary Perl program would. However, the C data structures of the preprocessor and parser are hidden behind the hooks and backcalls interfaces. All command line options accepted by **cpp** are also accepted by **PCp³** (this often makes it easier to use the tool in place of a compiler for analyzing complete packages as described by Murphy et al. [16]). Additionally, **PCp³** accepts a **--noparse** option that turns off its parser and the calls to related hooks. **PCp³** provides over currently over forty action hooks (see Appendix A).

The implementation of the main **PCp³** program, the (roughly thirty) backcalls (see Appendix B), and the glue connecting the components totals about 1,800 lines of C code. About 60% of this code deals directly with passing arguments between C and Perl.

Performance

The performance of **PCp³** depends on the complexity of the analysis. For comparison, for **gcc** to compile and optimize the 5,000 lines of the **bc** package [7] (an arbitrary precision arithmetic language interpreter) required 38 seconds on a Pentium Pro 200 Linux machine. Run-

ning **PCp³** with its test analysis consisting of 600 lines of hooks that exercise every event required 4 minutes for the **bc** source. Removing hooks for **CPP_OUT** and **TOKEN⁵** reduced the running time to 50 seconds. With all action code turned off, the running time is less than 10 seconds. The useful analyses described in this paper involve only a handful of callbacks, and thus execute very quickly.

5 RELATED WORK

Numerous tools exist for assisting the software engineer in understanding and manipulating source code. Griswold, Atkinson and McCurdy review a number of them while motivating their **TAWK** tool [9]. **TAWK** uses C as its action language and matches patterns in the abstract syntax tree. **TAWK**, like **PCp³**, operates on unprocessed source code. Instead of embedding a preprocessor, **TAWK** tries to parse each macro definition as an expression, allowing macro arguments to be types, as well. If that succeeds, the macro is left unexpanded in the code and becomes a node resembling a function call in their AST. About 92% of macro definitions in the two packages they studied parsed acceptably. For the remaining 8%, **TAWK** expands the macro uses before feeding the resulting tokens to their parser.

The Intentional Programming group at Microsoft Research [20], headed by Charles Simonyi, is interested in preserving preprocessor abstractions as they import legacy C code into their system. They developed a novel technique for handling preprocessor directives.⁶ Before preprocessing, conditional compilation directives are converted to stylized variable declarations. Then, the source text is preprocessed and all macros are expanded while marking each token with its textual derivation by the preprocessor. These declarations and the other source are then run through a C++ parser to create an AST. The annotations decorate the tree, and “enzymes” privity to the meaning of the stylized declarations process the tree in an attempt to identify abstractions. When macro expansions vary from use to use (e.g., **--LINE--**), the non-constant text is considered an extra argument to the macro, and the different expansions are explicitly passed at the invocation sites. Especially unusual uses of conditional compilation directives cause problems because of constraints on where C++ declarations may go, but generally the group is optimistic about the possibilities for their approach. Some of their techniques might

⁵These hooks are invoked on every string that **cpp** outputs and on every token it inputs, respectively. They are the most frequently invoked hooks.

⁶Charles Simonyi and Rammoorthy Sridharan, personal conversation.

be applicable to a future version of PCp³, especially as the AST-related backcalls mature.

LCLint [3] attempts to analyze macro definitions for dangerous or error-prone constructs. It allows the programmer to add annotations to the code in comments. These annotations give the LCLint checker extra information that it then verifies. For example, a macro argument can be marked that it needs to be side-effect free at an invocation site, and LCLint will generate a warning message if that constraint is violated. Evans’s focus is on finding errors, and dealing with macro expansions is largely ignored [3, Ch. 8]. Unlike PCp³, LCLint is not designed to be a general framework for analyses

Krone and Snelting analyze conditional compilation directives in the context of a lattice-theoretic framework for inferring configuration structures from the source code. They study how `#if` guards depend on and relate to each other, and provide a means of visualizing the relationships with the goal of improving the programmer’s understanding of the structure and properties of the configurations [12].

6 LIMITATIONS

Our framework provides a concise and flexible infrastructure for analyzing unprocessed C code. It permits reasoning about the entire source artifact and eliminates the need of individual analyses to mimic the preprocessor. Nevertheless there are several weaknesses of the framework and our implementation in PCp³.

First, some sophisticated preprocessor analyses (e.g., the second analysis in section 3) are dependent on the order that action hooks are called. This in turn requires an intimate understanding of the implementation’s pre-processing and (to a lesser extent) parsing peculiarities. Fortunately, the fast development time provided by the scripting language permits easy exploration and testing of analyses.

Because PCp³ works like a compiler, it handles only a single translation unit at a time, complicating whole-program analysis. The analyses discussed in section 3 output character-indexed annotations of the input that a separate tool later applies to the source code to perform the final transformation or to permit interactive exploration. These extra tools can also combine information derived from various translation units (e.g., marking expansions in a header file that is included by multiple source files). To better support reasoning about an entire source code artifact, a database approach similar to CIA++ [8] could be used. Atkinson and Griswold mention the importance of flexibility

in allowing the user to select the appropriate balance between precision and performance of a whole-program analysis [1]. One could provide this flexibility by using one of Perl’s data-persistence modules to permit specified data structures to be shared among invocations on separate translation units.

Better support for multiple conditional compilation versions would be useful. PCp³’s mechanisms for handling multiple source configurations are primitive—a single distinguished path (a version) through the source code is treated specially. Ideally, multiple versions could be considered with more uniform handling of the various paths. This would permit future blocks of code that are hidden via an `#ifdef` to be properly influenced by prior blocks of code using the same guard. Krone and Snelting suggest that the number of distinct paths through the source is reasonably bounded [12]. A generalized symbol table could track which configurations contain each symbol, and how the type of a variable depends on conditional compilation guards. A similar generalization could be made for the preprocessor name table.

Finally, although PCp³’s CTree-based parser constructs an abstract syntax tree, there is currently no easy way to access the AST from action callbacks. More backcalls and utility subroutines could be written to permit useful manipulation of the abstract syntax tree to avoid many of the problems created by the current limitation of a single pass over the source code. To make the AST more useful, some generalization of the tree could permit representation of preprocessor-specific annotations.

7 CONCLUSION

Our framework distinguishes itself from other software engineering tools by providing an accurate parse without disregarding the C preprocessor. By maintaining a tight mapping between the unprocessed and preprocessed code, analyses requiring expanded code can be considered in terms of the unprocessed code familiar to the programmer. This approach empowers and simplifies analyses while relieving the tool builder from the burden of partially re-implementing `cpp` for each desired analysis.

ACKNOWLEDGMENTS

This paper was supported by a National Science Foundation Graduate Fellowship. Any opinions, findings, conclusions, or recommendations expressed in this publication are those of the author, and do not necessarily reflect the views of the National Science Foundation.

Thanks to Michael Ernst, Craig Kaplan, Brian Michalowski, Gail Murphy, and Douglas Zongker for

their thoughtful comments on earlier revisions of this paper.

A ABRIDGED HOOKS REFERENCE

The below lists action hooks and the conditions under which the framework invokes the corresponding call-back. We omit partially redundant and less useful hooks. Passed parameters generally include source code character offsets and other relevant details of the invoking action.

STARTUP, STARTPARSE, EXIT

Initialization of `cpp`, initialization of parser, and conclusion of processing

HANDLE_DIRECTIVE

Reading of a `cpp` directive

DO_XIFDEF, DO_IF, DO_ELIF, DO_ELSE, DO_ENDIF

Handling of conditional compilation directives

CREATE_PREDEF, CREATE_DEF

Predefined (e.g., `__GCC__`) and user-defined macro definitions (from `#defines`)

DELETE_DEF

`#undef` of a macro

SPECIAL_SYMBOL

Expansion of a special symbol (e.g., `__FILE__`)

EXPAND_MACRO, MACRO_CLEANUP

In-code expansion of a macro, and completed expansion of a macro (these action hooks nest)

IFDEF_MACRO, IFDEF_LOOKUP_MACRO

Expansion of a macro inside a conditional compilation directive, or test for definedness

COMMENT, STRING_CONSTANT

Reading of a comment and string constant

CPP_ERROR, CPP_WARN, CPP_PEDWARN

`cpp` errors, warnings, and pedantic warnings

CPP_OUT, CPP_TOKEN

Outputting a sequence of characters written, and inputting a token

INCLUDE_FILE, DONE_INCLUDE_FILE

Inclusion of a file, and conclusion of reading an included file (these action hooks nest)

FUNCTION, FUNC_CALL

Parsing of a function definition and function invocation

ASSIGN_EXPR, EQUALITY_TEST

Parsing of an assignment and equality test

TYPEDDEF, VARDECL

Parsing of a typedef and variable declaration parsed.

POP_PERL_BUFFER

Completion of processing a Perl-pushed (via the `PushBuffer` backcall) buffer

B ABRIDGED BACKCALLS REFERENCE

The below backcalls are grouped by functionality. Parameters, if any, are indicated.

CbuffersBack, MacroExpansionHistory

Return the number of macro expansions deep in the current expansion, and the list of the stack of the current expansion history (the nesting of the various macro invocations).

FExpandingMacros

Return true if a macro is currently being expanded.

CchOffset, CcchOutput

Return a character position offset into the current input source file, and the character position offset into the output file.

InFname, Fname

Return the main filename (e.g., the one given on the command line), and the filename of the input currently being processed.

ExpansionLookup, LookupSymbol

Takes a macro or variable name; return the expansion of that macro or the symbol table entry of the identifier.

EnterScope, ExitScope

Push and pop new symbol table scopes.

PushHashTab, PopHashTab

Push and pop new macro definition hash table scopes.

ParseStateStack, SetStateStack

Return and set the stack of states for the parser (the latter takes a list of parser states).

YYGetState, YYSetState

Return and set the current state for the parser (the latter takes a parser state).

REFERENCES

- [1] D. C. Atkinson and W. G. Griswold. The design of whole-program analysis tools. In *Proceedings of the 18th International Conference on Software Engineering*, pages 16–27, March 1996.
- [2] M. Ernst, G. Badros, and D. Notkin. An empirical study of C preprocessor use. Technical Report UW-CSE-97-04-06, University of Washington, April 1997.
- [3] D. Evans. *LCLint User's Guide*. MIT Laboratory for Computer Science, Cambridge, MA, v2.2 edition, August 1996.
- [4] S. Flisakowski. CTree distribution, v0.5. Freely available software package, July 1997. ftp://ftp.kagi.com/flisakow/ctree_05.tar.gz.
- [5] F. S. Foundation. Bison v1.25 and flex v2.5 distributions. Freely available software package, April 1992,1995. <ftp://prep.ai.mit.edu/pub/>.
- [6] F. S. Foundation. GCC distribution, v2.7.2.2. Freely available software package, January 1997. <ftp://prep.ai.mit.edu/pub/gcc-2.7.2.tar.gz>.
- [7] F. S. Foundation. bc distribution. Freely available software package, June 1998. <ftp://prep.ai.mit.edu/pub/bc-1.05a.tar.gz>.
- [8] J. E. Grass and Y.-F. Chen. The C++ information abstractor. In *Proceedings of the USENIX 1992 C++ Conference*, Portland, Oregon, August 1992.
- [9] W. G. Griswold, D. C. Atkinson, and C. McCurdy. Fast, flexible syntactic pattern matching and processing. In *Proceedings of the IEEE 1996 Workshop on Program Comprehension*, March 1996.
- [10] S. P. Harbison. *C: A Reference Manual*. Prentice-Hall, Englewood Cliffs, New Jersey, 3rd edition, 1991.
- [11] B. W. Kernighan and D. M. Ritchie. *The C Programming Language*. Prentice-Hall, Englewood Cliffs, New Jersey, 2nd edition, 1988.
- [12] M. Krone and G. Snelting. On the inference of configuration structures from source code. In *Proceedings of the 16th International Conference on Software Engineering*, pages 49–57. IEEE Computer Society Press, May 1994.
- [13] J. R. Levine. *Lex & Yacc*. O'Reilly & Associates, Inc., Sebastopol, California, 2nd edition, 1992.
- [14] B. Lewis, D. LaLiberte, R. Stallman, and The GNU Manual Group. *GNU Emacs Lisp Reference Manual*. Free Software Foundation, Cambridge, Massachusetts, 2nd edition, May 1994.
- [15] G. C. Murphy and D. Notkin. Lightweight lexical source model extraction. *ACM Transactions on Software Engineering and Methodology*, 5(3):263–291, July 1996.
- [16] G. C. Murphy, D. Notkin, W. G. Griswold, and E. S. Lan. An empirical study of static call graph extractors. *ACM Transactions on Software Engineering and Methodology*, 7(2), April 1998.
- [17] J. Ousterhout. Scripting: Higher level programming for the 21st century. Web document, March 1997. <http://www.sunlabs.com/people/john.ousterhout>.
- [18] S. Reiss. *The Field Programming Environment: A Friendly Integrated Environment for Learning and Development*. Kluwer Academic Publishers, Hingham, MA, 1995.
- [19] M. Siff and T. Reps. Program generalization for software reuse: From C to C++. In *Sigsoft 96*. Sigsoft, 1996.
- [20] C. Simonyi. The intentional programming overview. Technical report, Microsoft Corporation, 1996. <http://www.research.microsoft.com/research/ip/main.htm>.
- [21] B. Stroustrup. *The Design and Evolution of C++*. Addison-Wesley, Reading, Massachusetts, 1994.